

# Week 3 – Introduction to Artificial Neural Networks

ELEC0144 Machine Learning for Robotics

Dr. Chow Yin Lai

Email: [uceecyl@ucl.ac.uk](mailto:uceecyl@ucl.ac.uk)

# Schedule

Week	Lecture	Workshop	Assignment Deadlines
1	Introduction; Image Processing	Image Processing	
2	Camera and Robot Calibration	Camera and Robot Calibration	
3	Introduction to Neural Networks	Camera and Robot Calibration	Friday: Camera and Robot Calibration
4	MLP and Backpropagation	MLP and Backpropagation	
5	CNN and Image Classification	MLP and Backpropagation	
6	Object Detection	MLP and Backpropagation	Friday: MLP and Backpropagation
7	Path Planning	Path Planning	
8	Kalman Filter SLAM	Path Planning	
9	Extended Kalman Filter SLAM	Path Planning	
10	Particle Filter SLAM	Path Planning	Friday: Path Planning

# Content

- Introduction to Neural Networks
- The Biological Neuron
- The Artificial Neuron
- Activation Functions
- Network Architectures
- Perceptrons
- Multilayer Perceptrons

# Content


- Introduction to Neural Networks
- The Biological Neuron
- The Artificial Neuron
- Activation Functions
- Network Architectures
- Perceptrons
- Multilayer Perceptrons

# What is a Neural Network (NN)? (1)

- A neural network is a **massively parallel distributed** processor that has a natural propensity for:
  - Storing experimental knowledge, and
  - Making it available for use.

# What is a Neural Network (NN)? (2)

- It is similar to the brain in two ways:
  - Knowledge is acquired by the network through a learning process.
  - Knowledge is stored using interneuron connection strengths known as synaptic weights.




## Biological vs Artificial NN

<https://www.quora.com/What-is-the-differences-between-artificial-neural-network-computer-science-and-biological-neural-network>

# Applications of NNs (1)

- NNs are mainly used for two types of applications:
- 1. Pattern Recognition or Classification
  - Example: Text recognition – Classify a handwritten alphabet as one of the 26 lower case letters.




# Applications of NNs (2)

- Another example: Differentiate between a cat and a dog.





[https://commons.wikimedia.org/wiki/File:Dog\\_Breeds.jpg](https://commons.wikimedia.org/wiki/File:Dog_Breeds.jpg)






# Applications of NNs (3)

- 2. Regression or Function Approximation
  - Example Left: Data fitting.
  - Example Right: House price prediction.



# Applications of NNs (4)

- Another example: To understand the input-output relationship of an inverted pendulum (Input = Force, Output = Angle).
- $\text{Angle} = \text{function}(\text{Force}) \rightarrow \text{Learnt by NN automatically!}$



# Learning of NNs (1)

- Recall that the knowledge is acquired through a **learning process**.
- For text recognition (e.g. “a”), we need to first show the NN many different handwritings of “a” – This is called **“Training”**.



## Learning of NNs (2)

- After training, we can test if NN recognizes a new (unseen before) “a” correctly – This is called “Testing / Generalisation”.
- This somewhat resembles human learning:
  - When you went to the kindergarten, you would see many different “a” from different teachers.
  - After some time, you learn to recognize the letter even if it was written by someone new.

## Learning of NNs (3)

- Similarly, for inverted pendulum (broom balancing), one would slowly **learn** the best hand motion in order to keep the broom upright.
- At the end, the brain would have learnt (unknowingly) the **relationship** between force and angular position!



<https://imgur.com/gallery/bZ919>

# Historical Perspective (1)

- 1943 McCulloch and Pitts published the **first description of an artificial NN**.
- 1950's – 60's First ANN developed by Marvin Minsky etc.: Single layer networks called **perceptron**, used for weather prediction, vision etc.
- 1970's Research virtually stopped, as many limitations were found e.g. incapable of solving many simple problems.

## Historical Perspective (2)


- 1982 Hopfield proposed associative memory model: NN evolves to minimize an energy function; renewed interest in NNs.
- 1986 **Back propagation** learning rule for **multi-layered networks** proposed, overcoming limitations of simple perceptrons. Explosion of research.
- 2010's **Deep neural network** achieved amazing results in vision and research in deep network explodes – It is now the state of the art for object recognition and many other applications.

# Content

- Introduction to Neural Networks
- **The Biological Neuron**
- The Artificial Neuron
- Activation Functions
- Network Architectures
- Perceptrons
- Multilayer Perceptrons




# One Biological Neuron



<https://www.quora.com/What-is-the-differences-between-artificial-neural-network-computer-science-and-biological-neural-network>)


# A Few Neurons Together




- Axon of one neuron almost touches dendrites of another neuron
- The small gap in between is the **Synapses**.
- Synapses can impose **excitation** (active) or **inhibition** (inactive) on the receptive neuron.

# Looking at the Neuron on the Right (1)

1. Input signals come from other neighbouring neurons.




# Looking at the Neuron on the Right (2)




2. The **synapses** impose **excitation** or **inhibition** of the signal. This can be thought of as multiplying with a **weight** of 0 or 1.

# Looking at the Neuron on the Right (3)

3. The cell body sums the incoming weighted signal.




# Looking at the Neuron on the Right (4)



4. When sufficient input is received (more than threshold), the neuron fires, i.e. generate a spike, which is transmitted to the axon.

# Looking at the Neuron on the Right (5)



5. If input is less than threshold, then no firing occurs.


# Content

- Introduction to Neural Networks
- The Biological Neuron
- **The Artificial Neuron**
- Activation Functions
- Network Architectures
- Perceptrons
- Multilayer Perceptrons




# Early Model of Neuron

- Very similar to biological neuron.




# Improved Model of Neuron

- Having threshold function not symmetrical about 0 makes some calculations difficult. Therefore, a bias is added.



# Further Improvement

- Various **activation functions** were also proposed in place of the threshold function.



# Mathematical Model of Neuron (1)


- As you can see, the model of a neuron has **three basic components**:
  - A set of **synapses** or connecting links, characterized by **weights**.
  - An **adder for summing** the weighted input signals.
  - An **activation function** to introduce nonlinearity into the output of a neuron, for e.g. limiting the amplitude of the output of a neuron within the range of  $[0, 1]$ ,  $[-1, 1]$ ,  $[0, \infty]$  etc.
    - The above ranges are due to “log sigmoid”, “tanh” and “rectified linear unit (ReLU)” respectively, which will be discussed later.

# Mathematical Model of Neuron (2)

- Mathematically, for a neuron k:

$$y_k = \varphi \left( \sum_{j=1}^m w_{jk} x_j + b_k \right)$$

$$y_k = \varphi(v_k)$$




# Mathematical Model of Neuron (3)

- Alternatively, we can think of the **bias** as input  $x_0 = 1$ , multiplied by weight  $w_{0k} = b_k$ .


$$y_k = \varphi \left( \underbrace{\sum_{j=0}^m w_{jk} x_j}_{v_k} \right)$$

$$y_k = \varphi(v_k)$$



# Mathematical Model of Neuron (4)

- It is useful to think of the neuron having **two halves**:
  - An adder  $\Sigma$ ,
  - Followed by activation function  $\varphi$ .




# Content

- Introduction to Neural Networks
- The Biological Neuron
- The Artificial Neuron
- **Activation Functions**
- Network Architectures
- Perceptrons
- Multilayer Perceptrons




# Activation Functions (1)

- Various activation functions have been proposed for neural networks.
- **Threshold function (hard-limiter):**
  - Note: McCulloch-Pits model (1943) of neuron used this form of function.




# Activation Functions (2)

- Piecewise linear function:
  - Linear combiner within certain a certain range, then saturated to 0 or 1.
  - If gradient of linear region is very high, it would reduce to threshold function.




# Activation Functions (3)

- **Sigmoid function** (s-shaped):
  - Commonly used in the past.
  - Strictly increasing function.
  - Asymptotically approach the saturation values.
- Continuous & Differentiable **everywhere** (very useful for optimisation later).
- Example: **Logistic function** (left) and **hyperbolic tangent function** (right).



$$\varphi(v) = \frac{1}{1 + e^{-av}}$$




$$\varphi(v) = \tanh(v)$$

# Activation Functions (4)

- Rectified Linear Unit (ReLU):

- Output = 0, if input is negative.
- Output = input, if input is positive.

- Mathematically: 
$$\varphi(v) = \max\{0, v\} = \begin{cases} 0, & v < 0 \\ v, & v \geq 0 \end{cases}$$




- Widely used nowadays, in the era of deep learning.
  - Multiplication of the gradient of sigmoid / tanh could become very small if there are many layers, leading to tiny change in weights and slow convergence (the “Vanishing Gradient” Problem)
  - On the other hand, the gradient of ReLU is 0 or 1, so after many layers the gradient will include the product of 1’s which is not too small.

# Activation Functions (5)

- Leaky ReLU:

- ReLU has a problem that if too many  $v$ 's are negative, then most of the ReLU output will simply be zero.
  - This would prohibit learning.
- Leaky ReLU solves this by giving a **small gradient when  $v$  is negative, for e.g.**



$$\phi(v) = \max\{0.01v, v\} = \begin{cases} 0.01v, & v < 0 \\ v, & v \geq 0 \end{cases}$$


# Activation Functions (6)

- Linear:

- The output is the same as the input.

$$\boxed{\varphi(v) = v}$$

- Although no nonlinearity is created by this unit, it is useful as the activation function for the **output layer of a multilayer NN**, particularly for regression problem.




## Example (1)

- A neuron  $k$  receives **inputs** from four other neurons whose activity levels are 10, -20, 6 and -2. The respective synaptic **weights** of neuron  $k$  are 0.7, 0.1, -1 and -0.4.
- Assume that the **bias** applied to the neuron is -10
- Calculate the output of neuron  $k$  for the following two situations:
  - a) The neuron is linear / ReLU / leaky ReLU.
  - b) The neuron is represented by a hard limiter.
  - c) The neuron is represented by Sigmoid / Tanh.

## Example (2)

- Solution: Let's calculate the weighted sum first.







## Example (3)

- Then calculate the output of activation function.

- Linear / ReLU / Leaky Relu:  $\varphi(1.8) = 1.8$




- Hard limiter:  $\varphi(1.8) = 1$




- Sigmoid (with  $a = 1$ ):  $\varphi(1.8) = \frac{1}{1+e^{-1.8}} = 0.8581$

- Tanh:  $\varphi(1.8) = \tanh(1.8) = 0.9468$




# Choice of Activation Functions (1)

- Later, you will learn about the terms “hidden layer” and “output layer”.
- For **hidden layer**, it is common to use ReLU, leaky ReLU, tanh and log sigmoid as activation functions.
- Usually selected based on trial-and-error, on a case-by-case basis.




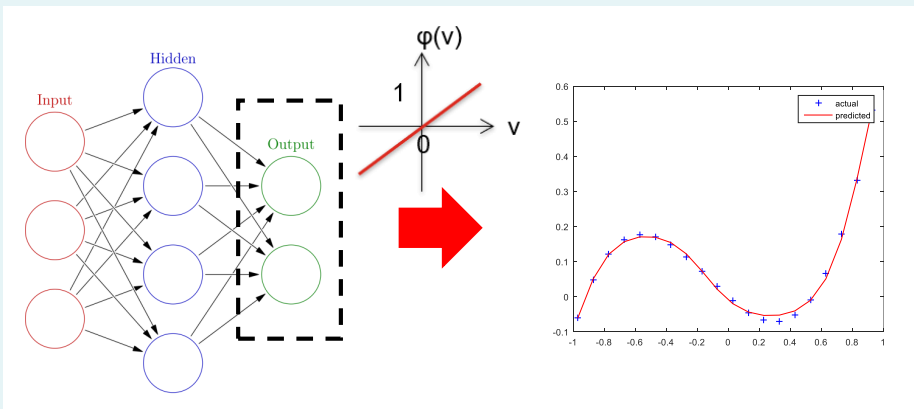
## Choice of Activation Functions (2)

- For **output layer**, the choice will depend on the application and data.
- For **classification**, in which the output is either 0 or 1, log sigmoid is an obvious choice.
  - E.g. output = 0.9  $\rightarrow$  90% confident that the image is a dog.



# Choice of Activation Functions (3)


- For **regression**, if the output is unconstrained (left), then linear function is an good option. It simply sums the nonlinearity created by the hidden layer.
- If the output is constrained (for e.g. right – house price cannot be less than 0!), then ReLU is a good choice.



# Differentiation of Activation Fcns (1)


- For training of the neural networks, the **weights will be updated** through **optimization** / minimization of cost function.
- This involves **differentiation of the cost function**.
- Let's look at the derivatives of several activation functions:
  - **ReLU:**

$$\varphi = \max\{0, v\} = \begin{cases} 0, & v < 0 \\ v, & v \geq 0 \end{cases} \rightarrow \frac{d\varphi}{dv} = \begin{cases} 0, & v < 0 \\ 1, & v \geq 0 \end{cases}$$



- **Leaky ReLU:**


$$\varphi(v) = \max\{0.01v, v\} = \begin{cases} 0.01v, & v < 0 \\ v, & v \geq 0 \end{cases} \rightarrow \frac{d\varphi}{dv} = \begin{cases} 0.01, & v < 0 \\ 1, & v \geq 0 \end{cases}$$



# Differentiation of Activation Fcns (2)

- Log Sigmoid:

$$\varphi(v) = \frac{1}{1+e^{-av}} = (1 + e^{-av})^{-1}$$



$$\begin{aligned}\frac{d\varphi}{dv} &= -(1 + e^{-av})^{-2} \cdot e^{-av} \cdot (-a) = a \cdot \frac{e^{-av}}{(1 + e^{-av})^2} \\ &= a \cdot \frac{e^{-av}}{1 + e^{-av}} \cdot \frac{1}{1 + e^{-av}} = a \cdot \frac{1 + e^{-av} - 1}{1 + e^{-av}} \cdot \frac{1}{1 + e^{-av}} \\ &= a \cdot \left(1 - \frac{1}{1 + e^{-av}}\right) \cdot \frac{1}{1 + e^{-av}} = a \cdot (1 - \varphi(v)) \cdot \varphi(v)\end{aligned}$$


$$\rightarrow \boxed{\frac{d\varphi}{dv} = a \cdot (1 - \varphi(v)) \cdot \varphi(v)}$$

- That means, if you have the output value  $\varphi(v)$ , then you can get the derivative easily!

# Differentiation of Activation Fcns (3)

- Tanh:

$$\boxed{\varphi(v) = \tanh(v) = \frac{e^v - e^{-v}}{e^v + e^{-v}}}$$



$$\begin{aligned} \frac{d\varphi}{dv} \text{ using quotient rule} &= \frac{(e^v + e^{-v})(e^v + e^{-v}) - (e^v - e^{-v})(e^v - e^{-v})}{(e^v + e^{-v})^2} \\ &= \frac{(e^v + e^{-v})^2 - (e^v - e^{-v})^2}{(e^v + e^{-v})^2} = 1 - \frac{(e^v - e^{-v})^2}{(e^v + e^{-v})^2} = 1 - \varphi^2(v) \end{aligned}$$

$$\rightarrow \boxed{\frac{d\varphi}{dv} = 1 - \varphi^2(v)}$$

- Again, if you have the output value  $\varphi(v)$ , then you can get the derivative easily!

# Content

- Introduction to Neural Networks
- The Biological Neuron
- The Artificial Neuron
- Activation Functions
- **Network Architectures**
- Perceptrons
- Multilayer Perceptrons




# Network Architectures (1)

- One single neuron is only able to solve some **very simple problems**.
- To solve for **more complex problems**, networks with large number of neurons are required.

## Network Architectures (2)

- Multilayer feedforward neural networks: Connections only from left to right.




- The above is a “5-3-3-2” network.

# Network Architectures (3)

- **Input Layer:**

- NO process carried out here. Only pass signal to the next layer.
- NOT a design parameter.
  - Number of nodes will be the same as **dimension** of inputs.



# Network Architectures (4)

- Hidden Layer:
  - Design parameters:
    - Number of hidden layers.
    - Number of nodes in each hidden layer.
    - Activation functions – can be different for different layers.
  - It can be mathematically proven that **one hidden layer is enough** to approximate any bounded continuous function.
    - So why add more hidden layers?


# Network Architectures (5)

- **Advantage** of having more hidden layers:
  - The **total number of synapses weight** in multilayer network is less → less parameters to tune
    - E.g. 1-9-1 network: 28 weights (including biases)
    - 1-3-3-1 network: 22 weights (including biases)
- **Disadvantage** of having more hidden layers:
  - More prone to **local minima** due to its more complicated structure.

# Network Architectures (6)

- Output Layer:

- Number of nodes is NOT a design parameters. Will be the same as dimension of outputs (# functions for regression, # classes for classification).
- Activation function is a design parameter.
- Depends on the expected output, as already discussed in the section “activation function”.



# Content

- Introduction to Neural Networks
- The Biological Neuron
- The Artificial Neuron
- Activation Functions
- Network Architectures
- **Perceptrons**
- Multilayer Perceptrons


# Perceptron (1)

- Earlier on, we mentioned that knowledge is acquired by a neural network **through learning**.
- Let's use a simple network – a “**perceptron**” – to demonstrate what this means.



## Perceptron (2)

- Perceptron is the simplest form of a NN for classification of patterns.
- It learns via examples, how to **assign input vectors** (samples) to different **classes** (Rosenblatt, 1958).




## 2D example (1)

- A 2-dimensional example: To correctly classify the external inputs  $\{x_1, x_2\}$  into one of two classes  $\{C_1$  or  $C_0\}$ .  
2D
- E.g. AND problem:

X1 (input)	0	0	1	1
X2 (input)	0	1	0	1
Y (output)	0	0	0	1


## 2D example (2)

- The perceptron is shown below. We want to find  $b$ ,  $w_{k1}$  and  $w_{k2}$ , so that when given  $x_1$  and  $x_2$ , the correct  $y$  will be calculated.




# Analytical Solution (1)

- Let's solve this analytically first.
- This is a simple **2-dimensional** problem, thus we can sketch the input-output space:



## Analytical Solution (2)

- The classes are **linearly separable** and we can easily get a straight line to separate the two classes.




## Analytical Solution (3)

- From the line equation, we would either have:
  - $v_k = -x_1 - x_2 + 1.5$ , or
  - $v_k = x_1 + x_2 - 1.5$ .
- Which one is correct?
- Try with  $x_1 = 0, x_2 = 0$ :
  - First equation gives  $v_k = 1.5$ , then  $\varphi(1.5) = 1 \rightarrow$  Wrong
  - Second equation gives  $v_k = -1.5$ , then  $\varphi(-1.5) = 0 \rightarrow$  Correct!

# Analytical Solution (4)

- Thus, the complete perceptron is as follows:
  - (from  $v_k = x_1 + x_2 - 1.5$ )



# Learning Method (1)

- We will now use a learning procedure to train the perceptron.
  - A training set of input-output vectors, i.e. exemplars, is given.
  - The weight vector will be tuned in such a way that the best classification of the training vectors is achieved.
- The learning procedure works even for more general cases (more than 2D), so we will define some terms on the next page.



# Learning Method (2)

- The **input vector** is:

$$x'(t) = [1, x_1(t), x_2(t), \dots, x_m(t)]^T$$

- The **weight vector** is:

$$w'(t) = [b(t), w_1(t), w_2(t), \dots, w_m(t)]^T$$

- Where  $t$  denotes the **iteration step**.
- The **intermediate output** (before the hard limiter) is:

$$v(t) = w'^T(t) \cdot x'(t)$$

## Learning Method (3)

- The equation  $w'^T \cdot x' = 0$ , plotted in an  $m$ -dimensional space with coordinates  $x_1, x_2, \dots, x_m$ , defines a **hyperplane** which separates the two classes.
- With the help of hard limiter, we finally get:

$$w'^T \cdot x' \geq 0 \text{ Class 1}$$

$$w'^T \cdot x' < 0 \text{ Class 0}$$

# Perceptron Learning Algorithm

- Start with randomly chosen weight vector  $w'(0)$ .
- Let  $t = 1$ .
- **While** there exist input vectors that are wrongly classified by  $w'(t - 1)$ , **do**
  - If  $x'$  is a misclassified input vector,
  - Update the weight vector to

$$w'(t) = w'(t - 1) + \eta(d - y)x'$$

- Where  $\eta > 0$  and  $d = \begin{cases} 1 & \text{if } x \text{ belongs to Class 1} \\ 0 & \text{if } x \text{ belongs to Class 0} \end{cases}$ , and  $d = \text{desired output}$   
 $y = \text{network output}$
- Increment  $n$
- **End While**

# AND-Example Revisited (1-1)

X1 (input)	0	0	1	1
X2 (input)	0	1	0	1
d (desired output)	0	0	0	1

- Randomly initiated weight:  $w'(0) = [0.5, 0.5, 0.5]^T$

- First epoch, first column of data

$$w'(0)^T \cdot x = [0.5, 0.5, 0.5] \cdot \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = 0.5 \rightarrow y = \varphi(0.5) = 1, d = 0 \rightarrow \text{misclassified}$$

- Because of misclassification, update  $w$  to:

$$w'(1) = w'(0) + \underbrace{\eta}_{0.1} (d - y)x' = \begin{bmatrix} 0.5 \\ 0.5 \\ 0.5 \end{bmatrix} + 0.1(0 - 1) \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} 0.4 \\ 0.5 \\ 0.5 \end{bmatrix} \rightarrow \text{(Goes to next step)}$$

# AND-Example Revisited (1-2)

X1 (input)	0	0	1	1
X2 (input)	0	1	0	1
d (desired output)	0	0	0	1

- After previous step:  $w'(1) = [0.4, 0.5, 0.5]^T$

- First epoch, second column of data

$$w'(1)^T \cdot x = [0.4, 0.5, 0.5] \cdot \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix} = 0.9 \rightarrow y = \varphi(0.9) = 1, d = 0 \rightarrow \text{misclassified}$$

- Because of misclassification, update  $w$  to:

$$w'(2) = w'(1) + \eta(d - y)x' = \begin{bmatrix} 0.4 \\ 0.5 \\ 0.5 \end{bmatrix} + 0.1(0 - 1) \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix} = \begin{bmatrix} 0.3 \\ 0.5 \\ 0.4 \end{bmatrix}$$

(Goes to  
next step)



# AND-Example Revisited (1-3)

X1 (input)	0	0	1	1
X2 (input)	0	1	0	1
d (desired output)	0	0	0	1

- After previous step:  $w'(2) = [0.3, 0.5, 0.4]^T$

- First epoch, third column of data

$$w'(2)^T \cdot x = [0.3, 0.5, 0.4] \cdot \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} = 0.8 \rightarrow y = \varphi(0.8) = 1, d = 0 \rightarrow \text{misclassified}$$

- Because of misclassification, update  $w$  to:

$$w'(3) = w'(2) + \eta(d - y)x' = \begin{bmatrix} 0.3 \\ 0.5 \\ 0.4 \end{bmatrix} + 0.1(0 - 1) \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} = \begin{bmatrix} 0.2 \\ 0.4 \\ 0.4 \end{bmatrix}$$

(Goes to  
next step)



# AND-Example Revisited (1-4)

X1 (input)	0	0	1	1
X2 (input)	0	1	0	1
d (desired output)	0	0	0	1

- After previous step:  $w'(3) = [0.2, 0.4, 0.4]^T$

- First epoch, fourth column of data

$$w'(3)^T \cdot x = [0.2, 0.4, 0.4] \cdot \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} = 1 \rightarrow y = \varphi(1) = 1, d = 1 \rightarrow \text{correct}$$

- Because the classification is correct,  $w$  remains the same.

$$w'(4) = w'(3) = [0.2, 0.4, 0.4]^T$$

- We have **completed the first epoch**, i.e. all the data has been presented once.

## AND-Example Revisited (2-1)

X1 (input)	0	0	1	1
X2 (input)	0	1	0	1
d (desired output)	0	0	0	1

- After previous step:  $w'(4) = [0.2, 0.4, 0.4]^T$
- Continue to the second epoch, first column of data

$$w'(4)^T \cdot x = [0.2, 0.4, 0.4] \cdot \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = 0.2 \rightarrow y = \varphi(0.2) = 1, d = 0 \rightarrow \text{misclassified}$$

- Because of misclassification, update  $w$  to:

$$w'(5) = w'(4) + \eta(d - y)x' = \begin{bmatrix} 0.2 \\ 0.4 \\ 0.4 \end{bmatrix} + 0.1(0 - 1) \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} 0.1 \\ 0.4 \\ 0.4 \end{bmatrix}$$

(Goes to  
next step)



## AND-Example Revisited (2-2)

X1 (input)	0	0	1	1
X2 (input)	0	1	0	1
d (desired output)	0	0	0	1

- After previous step:  $w'(5) = [0.1, 0.4, 0.4]^T$

- Second epoch, second column of data

$$w'(5)^T \cdot x = [0.1, 0.4, 0.4] \cdot \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix} = 0.5 \rightarrow y = \varphi(0.5) = 1, d = 0 \rightarrow \text{misclassified}$$

- Because of misclassification, update  $w$  to:

$$w'(6) = w'(5) + \eta(d - y)x' = \begin{bmatrix} 0.1 \\ 0.4 \\ 0.4 \end{bmatrix} + 0.1(0 - 1) \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix} = \begin{bmatrix} 0 \\ 0.4 \\ 0.3 \end{bmatrix}$$

(Goes to  
next step)



## AND-Example Revisited (2-3)

X1 (input)	0	0	1	1
X2 (input)	0	1	0	1
d (desired output)	0	0	0	1


- After previous step:  $w'(6) = [0, 0.4, 0.3]^T$

- Second epoch, third column of data

$$w'(6)^T \cdot x = [0, 0.4, 0.3] \cdot \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} = 0.4 \rightarrow y = \varphi(0.4) = 1, d = 0 \rightarrow \text{misclassified}$$

- Because of misclassification, update  $w$  to:

$$w'(7) = w'(6) + \eta(d - y)x' = \begin{bmatrix} 0 \\ 0.4 \\ 0.3 \end{bmatrix} + 0.1(0 - 1) \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} = \begin{bmatrix} -0.1 \\ 0.3 \\ 0.3 \end{bmatrix}$$

(Goes to next step) 

## AND-Example Revisited (2-4)

X1 (input)	0	0	1	1
X2 (input)	0	1	0	1
d (desired output)	0	0	0	1

- After previous step:  $w'(7) = [-0.1, 0.3, 0.3]^T$

- Second epoch, fourth column of data

$$w'(7)^T \cdot x = [-0.1, 0.3, 0.3] \cdot \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} = 0.5 \rightarrow y = \varphi(0.5) = 1, d = 1 \rightarrow \text{correct}$$

- Because the classification is correct,  $w$  remains the same.

$$w'(8) = w'(7) = [-0.1, 0.3, 0.3]^T$$

- Second epoch completed.

# AND-Example Revisited (3-1)

X1 (input)	0	0	1	1
X2 (input)	0	1	0	1
d (desired output)	0	0	0	1

- After previous step:  $w'(8) = [-0.1, 0.3, 0.3]^T$

- Continue to **third epoch, first column of data**

$$w'(8)^T \cdot x = [-0.1, 0.3, 0.3] \cdot \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = -0.1 \rightarrow y = \varphi(-0.1) = 0, d = 0 \rightarrow \text{correct}$$

- Because the classification is correct,  $w$  remains the same.

$$w'(9) = w'(8) = \begin{bmatrix} -0.1 \\ 0.3 \\ 0.3 \end{bmatrix}$$

(Goes to  
next step)



## AND-Example Revisited (3-2)

X1 (input)	0	0	1	1
X2 (input)	0	1	0	1
d (desired output)	0	0	0	1

- After previous step:  $w'(9) = [-0.1, 0.3, 0.3]^T$

- Third epoch, second column of data

$$w'(9)^T \cdot x = [-0.1, 0.3, 0.3] \cdot \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix} = 0.2 \rightarrow y = \varphi(0.2) = 1, d = 0 \rightarrow \text{misclassified}$$

- Because of misclassification, update  $w$  to:

$$w'(10) = w'(9) + \eta(d - y)x' = \begin{bmatrix} -0.1 \\ 0.3 \\ 0.3 \end{bmatrix} + 0.1(0 - 1) \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix} = \begin{bmatrix} -0.2 \\ 0.3 \\ 0.2 \end{bmatrix} \rightarrow \text{(Goes to next step)}$$

## AND-Example Revisited (3-3)

X1 (input)	0	0	1	1
X2 (input)	0	1	0	1
d (desired output)	0	0	0	1

- After previous step:  $w'(10) = [-0.2, 0.3, 0.2]^T$

- Third epoch, third column of data

$$w'(10)^T \cdot x = [-0.2, 0.3, 0.2] \cdot \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} = 0.1 \rightarrow y = \varphi(0.1) = 1, d = 0 \rightarrow \text{misclassified}$$

- Because of misclassification, update  $w$  to:

$$w'(11) = w'(10) + \eta(d - y)x' = \begin{bmatrix} -0.2 \\ 0.3 \\ 0.2 \end{bmatrix} + 0.1(0 - 1) \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} = \begin{bmatrix} -0.3 \\ 0.2 \\ 0.2 \end{bmatrix} \rightarrow \text{(Goes to next step)}$$


## AND-Example Revisited (3-4)

X1 (input)	0	0	1	1
X2 (input)	0	1	0	1
d (desired output)	0	0	0	1

- After previous step:  $w'(11) = [-0.3, 0.2, 0.2]^T$
- We can continue doing the same, and it will be observed that with  $w' = [-0.3, 0.2, 0.2]^T$  always gives the **correct classification**!
- The perceptron is successfully trained after 3 epochs!

# AND-Example Revisited (4)

- **Summary:** After three epochs, we arrive at the weight  $w' = [-0.3, 0.2, 0.2]^T$  which correctly classifies all the data point.
- It *looks* different from what we manually calculated ( $w'_{\text{manual}} = [-1.5, 1, 1]^T$ ) but is actually equivalent!
  - Manual:  $x_2 = -x_1 + 1.5$
  - Trained:  $0.2 x_2 = -0.2 x_1 + 0.3$





# MATLAB Code (1)

```
clear all
close all
clc

#####
% Data %
#####

% In the order of [x0 x1 x2 d]
Data = [1 0 0 0;
        1 0 1 0;
        1 1 0 0;
        1 1 1 1];

#####
% Parameters %
#####

w = [0.5, 0.5, 0.5]'; % [bias w1 w2]
eta = 0.1; % Try changing this and observe results
epochs = 4;
wrecord = w;
```

```
#####
% Algorithm %
#####


[ndata, mdata] = size(Data);

for i = 1:epochs
    for j = 1:ndata
        x = Data(j, 1:3)';
        v = w'*x;
        if v >= 0
            y = 1;
        else
            y = 0;
        end
        d = Data(j, 4);
        w = w + eta*(d-y)*x;
        wrecord = [wrecord w];
    end
end

figure, plot(wrecord(1,:))
hold on, plot(wrecord(2,:))
hold on, plot(wrecord(3,:))
legend('bias', 'w1', 'w2')
```

# MATLAB Code (2)

- Progress of the weights during training.



# Learning Rate (1)


- The parameter  $\eta > 0$  influences the learning rate.
- Small value leads to slow learning.
- Large value can “spoil” the learning that has taken place earlier with respect to other data points.
- Therefore, some medium value is the best.
  - What “medium” means depends on the problem being solved.

## Learning Rate (2)

- **Exercise:** Try changing the  $\eta$  value in the MATLAB code earlier, and observe the results.
  - Note: You might need to increase the epochs to allow the weights to converge.


# Limitations of Perceptrons (1)

- Perceptrons can only classify two classes which are **linearly separable**:
  - E.g. **2-D case**, if the classes cannot be separated by a **straight line**, then they cannot be classified by simple perceptrons.




# Limitations of Perceptrons (2)

- **AND** – Linearly separable, as already seen earlier.
- **OR** – also linearly separable:



- **XOR** - ???




# Content

- Introduction to Neural Networks
- The Biological Neuron
- The Artificial Neuron
- Activation Functions
- Network Architectures
- Perceptrons
- **Multilayer Perceptrons**

# XOR Problem (1)


- We saw that a single perceptron is not able to solve the XOR problem.






# XOR Problem (2)

- To do that, we need **a few perceptrons** working together.
- For instance, consider the following network by Touretzky and Pomerleau (1989):




# XOR Problem (3)

- Let's check the results:
- Case 1:  $x_1 = 0, x_2 = 0$ :




# XOR Problem (4)

- Case 2:  $x_1 = 1, x_2 = 0$ :




# XOR Problem (5)

- Case 3:  $x_1 = 0, x_2 = 1$ :



# XOR Problem (6)


- Case 4:  $x_1 = 1, x_2 = 1$ :



Correctly  
solves  
the XOR  
problem!


# Multilayer Perceptrons MLP (1)

- As can be seen, by connecting multiple perceptrons together, more complicated problems can be solved.
- A **Multilayer Perceptron (MLP)** consists of:
  - An **input layer** (note: no computation here).
  - One or more **hidden layers** of computation nodes.
  - An **output layer** of computation nodes.
- E.g.



# Multilayer Perceptrons MLP (2)

- Again, it is useful to think of the hidden layer nodes and output layer nodes as **two halves** (adder + activation):



- For convenience, we will also not explicitly draw out the **biases**.

# Multilayer Perceptrons MLP (3)

- As mentioned previously, MLP generally uses **differentiable activation functions** instead of threshold function.
- This is because for training of the weights, we perform **optimization** which require the functions to be differentiable.



**Thank you for your attention!**

Any questions?