

# **COMP0104 Software Development Practice: Profiling and Coverage**

**Jens Krinke**

Centre for Research on Evolution, Search & Testing  
Software Systems Engineering Group  
Department of Computer Science  
University College London

# Overview

- Tuning: Improving the Performance
- Profiling: Create Profiles over Program Execution
- Coverage: Measure the Program Execution

# Improving performance

- A important part of programming is **tuning**, the systematic performance improvement.
- Systematic does not mean randomly!
- With profiling, one can determine
  - the **locations** where the performance can be improved
  - the **effect** of possible optimisations on the runtime.

# System Tuning



Establish acceptable behaviour



Measure the current performance – is it acceptable?



Identify the bottleneck(s)



Modify the bottleneck



Measure the performance again – is it now acceptable?

# Tools for profiling

- TIME determines the runtime of **processes**.
- GPROF shows how the runtime is split across **functions**.
- GCOV counts how often individual **program lines** are executed.

# Runtime measurement using TIME

TIME executes the command passed as an argument, and then outputs the runtime that was used by the process.

```
$ c++ -o huffencode huffencode.C  
$ time ./huffencode COPYING > COPYING.huff.C  
real      0m0.465s  
user      0m0.420s  
sys       0m0.010s  
$ _
```

# Runtime optimisations

```
$ c++ -O -o huffencode huffencode.C  
$ time ./huffencode COPYING > COPYING.huff.C  
real      0m0.735s  
user      0m0.130s  
sys       0m0.010s  
$ _
```

CPU time for user functions has dropped to a third

However, the real runtime has increased. Why?

- other processes could be running
- it may take a moment for the newly created program to be loaded into memory

# Program profiles with GPROF

A program profile shows

- which functions have been called,
- how often, and
- how much time they required.

This is done in three steps:

- During compilation, the program is instrumented
- During execution, a program profile is created.
- The program profile is analysed by GPROF.



# Program profiles with GPROF

- Set up the program for profiling:

```
$ c++ -O -pg -o huffencode huffencode.C  
$ _
```

- Execute the program

```
$ ./huffencode COPYING > COPYING.huff.C  
$ _
```

- Analysis of the profile:

```
$ gprof ./huffencode  
...
```

# The output from GPROF consists of two parts

- A **flat profile**, which indicates how the runtime is divided into individual functions
- A **structured profile** in which the functions that have been called and the functions that will be called are indicated

# The flat profile

% cumulative time	seconds	self seconds	calls	self us/call	total us/call	name
33.33	0.02	0.02	47110	0.42	0.42	string_Scat
16.67	0.03	0.01	10674	0.94	0.94	bits_to_byte
16.67	0.04	0.01	10673	0.94	0.94	string::at
16.67	0.05	0.01	1	10000.00	42298.03	encode
16.67	0.06	0.01	1	10000.00	17668.44	read_input
0.00	0.06	0.00	10983	0.00	0.00	string_Salloc
0.00	0.06	0.00	256	0.00	0.00	compare
0.00	0.06	0.00	154	0.00	0.00	string::op char *
0.00	0.06	0.00	154	0.00	0.44	_cook
0.00	0.06	0.00	153	0.00	0.00	extract_min
0.00	0.06	0.00	153	0.00	0.00	insert
0.00	0.06	0.00	78	0.00	0.00	string_Scopy
0.00	0.06	0.00	1	0.00	0.00	global destructors
0.00	0.06	0.00	1	0.00	0.00	global constructors
0.00	0.06	0.00	1	0.00	0.00	string::from
0.00	0.06	0.00	1	0.00	0.00	huffman
0.00	0.06	0.00	1	0.00	64.53	init_codes
0.00	0.06	0.00	1	0.00	0.00	initial_queue
0.00	0.06	0.00	1	0.00	0.00	length

# Insights

- The `string_Scat` function takes up a third of the runtime itself and with more than 47,000 calls is also the most frequently called function.
- Optimisations in the following functions can also increase the performance of the program:
  - `bits_to_byte`,
  - `string::at`,
  - `encode` **and**
  - `read_input`.

# Structured profiles

For each individual function  $f$ , a structured profile indicates

- what share  $f$  and the functions called by  $f$  have of the entire runtime (`%time`).
- which functions have called  $f$  and how often, as well as  $f$ 's individual share of their runtime (these functions are listed above  $f$ )
- which functions  $f$  has called and how often, and what share of  $f$ 's runtime they have (these functions are listed below  $f$ )

index	% time	self	children	called	name
		0.00	0.06	1/1	_start [2]
[1]	100.0	0.00	0.06	1	<b>main</b> [1]
		0.01	0.03	1/1	encode [3]
		0.01	0.01	1/1	read_input [5]
		0.00	0.00	1/1	write_huffman [10]
		0.00	0.00	1/1	huffman [17]
		0.00	0.00	1/1	write_encoding [20]
-----					
					<spontaneous>
[2]	100.0	0.00	0.06		<b>_start</b> [2]
		0.00	0.06	1/1	main [1]
-----					
		0.01	0.03	1/1	main [1]
[3]	70.5	0.01	0.03	1	<b>encode</b> [3]
		0.01	0.00	28737/47110	string_Scat [4]
		0.01	0.00	10674/10674	bits_to_byte [6]
		0.01	0.00	10673/10673	string::at [7]
		0.00	0.00	1/1	init_codes [9]
		0.00	0.00	77/154	_cook [8]
		0.00	0.00	10752/10983	string_Salloc [11]
		0.00	0.00	256/256	compare [12]
		0.00	0.00	1/1	string::from [16]
		0.00	0.00	1/78	string_Scopy [15]

		0.00	0.00	152/47110	init_codes [9]
		0.00	0.00	158/47110	_cook [8]
		0.01	0.00	18063/47110	read_input [5]
		0.01	0.00	28737/47110	encode [3]
[4]	33.3	0.02	0.00	47110	<b>string_Scat</b> [4]
-----					
		0.01	0.01	1/1	main [1]
[5]	29.4	0.01	0.01	1	<b>read_input</b> [5]
		0.01	0.00	18063/47110	string_Scat
[4]					
-----					

- `string_Scat` does not call any other functions (at least none that is set up for profiling)
- of the 47,110 calls,
  - 18,063 come from the `read_input` function
  - and 28,737 come from the `encode` function.

Perhaps these are two suitable candidates for optimisation?

# Coverage measurement

- GCOV measures the **line coverage**:  
it calculates how often each line of the program was run.
- Coverage is a measurement of how good a test suite is.
- Line coverage is useful for performance analyses.



# Setting up a program for coverage measurement

- During compilation, the program is set up for coverage measurement:

```
$ c++ -g -fprofile-arcs -ftest-coverage -o  
huffencode huffencode.C
```

- During the execution of the program, coverage measurement information is generated.

```
$ ./huffencode COPYING > COPYING.huff.C
```

- The coverage information is analyzed by GCOV

```
$ gcov huffencode.C
```

# GCOV output

```
// Fill CODES[C] with a [01]+ string denoting
// the encoding of C in TREE
static void init_codes(string codes[ UCHAR_MAX + 1],
                      HuffNode *tree,
                      string prefix = "")

153 {
153     if (tree == 0)
#####         return;

153     if (tree->isleaf)
77         codes[(unsigned char)tree->l.c] = prefix;
    else
76     {
76         init_codes(codes, tree->i.left,  prefix + "0");
76         init_codes(codes, tree->i.right, prefix + "1");
    }
}
```

```
// Encode TEXT using the Huffman tree TREE
static string encode(const string& text, HuffNode *tree)
1 {
1     string codes[ UCHAR_MAX + 1];

1     init_codes(codes, tree);
1     string bit_encoding;
1     for (int i = 0; i < text.length(); i++)
18063         bit_encoding += codes[(unsigned char)text[i]];

1     string byte_encoding;
1     for (int i = 0; i < bit_encoding.length() - BITS_PER_CHAR;
        i += BITS_PER_CHAR)
    {
10673         byte_encoding += bits_to_byte(bit_encoding.at(i, BITS_PER_CHAR));
10673     }
1     byte_encoding += bits_to_byte(bit_encoding.from(i));

1     return byte_encoding;
    }
```

# Insights

Now we can see from where the numerous calls of `string_Scat` come from:

- In two loops, character strings are concatenated.
- The first loop is executed once for every character of the input, precisely 18,063 times
- the second for each character of the output, precisely 10,673 times.

These two critical positions can now be optimised.

```
// Encode TEXT using the Huffman tree TREE
static string encode(const string& text, HuffNode *tree)
{
    string codes[ UCHAR_MAX + 1];

    init_codes(codes, tree);
    ostream bit_encoding_os;
    for (int i = 0; i < int(text.length()); i++)
        bit_encoding_os << codes[(unsigned char)text[i]];
    string bit_encoding(bit_encoding_os);

    ostream byte_encoding;
    for (int i = 0; i < int(bit_encoding.length()) - BITS_PER_CHAR;
         i += BITS_PER_CHAR)
    {
        byte_encoding << bits_to_byte(bit_encoding.at(i, BITS_PER_CHAR));
    }
    byte_encoding << bits_to_byte(bit_encoding.from(i));

    return byte_encoding;
}
```

# How GPROF and GCOV work

- GPROF and GCOV have the compiler to produce **instrumented code**, the code is enriched with specific instructions for profiling.
- For GPROF, functions are compiled so that every called function creates an entry showing when and from where it was called.
- For GCOV, each node in the control flow graph is provided with a counter that is incremented during the execution of the program.

# Sampling vs. instrumentation

- The run-time profiling of GPROF is not realised via counting, but via **sampling**.
- Instead, a special function analyses the current program counter at specific intervals and marks which functions have just been executed.  
(For example, 100 times per CPU second.)
- Compared to the accurate counting of function calls, sampling is subject to statistical inaccuracies.

# Testing with Coverage

- The data gained in the profiling can also be used for systematic tests.
- Through repeated execution of the program one can strive for a minimum coverage.
- Example:  
one can ensure that a test executes 95% of all statements at least once, thus making a statement about the quality of the test suite.



# Statement Coverage

- Statement coverage:  
% of statements executed by your testing
- Statement coverage =  $\frac{\text{\textit{\# of executed statements}}}{\text{\textit{\# of all statements}}}$
- Variations:
  - Instruction coverage
  - Line coverage

# EclEmma JaCoCo

The screenshot displays the Eclipse IDE interface. The left sidebar shows the 'JUnit' view with a green bar indicating successful test results. The main editor shows the source code of 'CloneReader.java' with lines 1-29. The right sidebar shows the 'Run' menu with options for coverage and breakpoints.

**JUnit View (Left Sidebar):**

- Finished after 0.07 seconds
- Runs: ✖ Errors: ✖ Failures: ✔
- RationalTest [Runner: JUnit 4] (C)
  - testEquality (0.001 s)
  - testNonEquality (0.000 s)

**Source Code (Main Editor):**

```

1  import java.io.*; /** Sample */
2
3  class Rational {
4      long numerator, denominator;
5
6      class Illegal extends Exception {
7          String reason;
8          Illegal (String reason) {
9              this.reason = reason;
10         }
11     }
12
13     Rational ()
14     {
15         numerator = 0;
16         denominator = 1;
17     }
18
19     Rational(long numerator, long denominator)
20     {
21         long num;
22         if (denominator == 0)
23             throw new Illegal ("Error, denominator is 0");
24         else if (denominator < 0)
25         {
26             denominator *=-1;
27             numerator *=-1;
28         }
29     }

```

**Run Menu (Right Sidebar):**

- Coverage Last Launched ^⌘F11
- Run ⇧⌘F11
- Debug ⌘F11
- Run History ▶
- Run As ▶
- Run Configurations...
- Debug History ▶
- Debug As ▶
- Debug Configurations...
- Coverage History ▶
- Coverage As ▶
- Coverage...
- Toggle Breakpoint ⇧⌘B
- Toggle Line Breakpoint
- Toggle Method Breakpoint
- Toggle Watchpoint
- Skip All Breakpoints
- Remove All Breakpoints
- Add Java Exception Breakpoint...
- Add Class Load Breakpoint...
- All References...
- All Instances... ⇧⌘N
- Instance Count

# EclEmma – Coverage View

RationalTest (Feb 28, 2011 11:32:34 AM)				
Element	Coverage	Covered Instructions	Missed Instructions ▼	Total Instructions
▼ src	48.5 %	167	177	344
▼ (default package)	48.5 %	167	177	344
▼ Rational.java	40.1 %	112	167	279
▼ Rational	40.1 %	112	167	279
● compareTo(Object)	0.0 %	0	53	53
● add(Rational)	0.0 %	0	23	23
● subtract(Rational)	0.0 %	0	23	23
● divide(Rational)	0.0 %	0	17	17
● multiply(Rational)	0.0 %	0	17	17
▲ Rational(long, long)	57.6 %	19	14	33
▶ Rational	0.0 %	0	9	9
▲ Rational()	0.0 %	0	9	9
● toString()	90.0 %	18	2	20
● equals(Object)	100.0 %	39	0	39
■ GCD(long, long)	100.0 %	15	0	15
■ simplestForm()	100.0 %	21	0	21

# Statement and branch coverage

- **Statement coverage** is only a minimal criterion.
- In practice, it is usually required that each branch of the control flow graph is run at least once (so-called **branch coverage**).
- Branch coverage = 
$$\frac{\text{\textit{\# of executed branching edges}}}{\text{\textit{\# of all branching edges}}}$$
- GCOV has the `-b` option which can be used to measure branch coverage.

# Branch coverage with GCOV

```
$ gcov -b huffencode.C
```

```
96.69% of 121 source lines executed
```

```
75.00% of 100 branches executed
```

```
69.00% of 100 branches taken at least once
```

```
92.36% of 144 calls executed
```

```
Creating huffencode.C.gcov.
```

```
$ _
```

```
153 {  
153     if (tree == 0)  
branch 0 taken = 100%  
#####         return;  
branch 0 never executed  
  
153     if (tree->isleaf)  
branch 0 taken = 50%  
77         {  
77             codes[(unsigned char)tree->l.c] = prefix;  
branch 1 taken = 100%  
             }  
            else  
76         {  
76             init_codes(codes, tree->i.left,  prefix + "0");  
76             init_codes(codes, tree->i.right, prefix + "1");  
branch 2 taken = 100%  
            }  
        }
```

# Insights

- The branch of the term `tree == 0` was executed in 100% of all runs.
- The branch from `return` to the end of the function was never executed (since the `return` statement was never executed).
- The branch of the term `tree->isleaf` (branch 0) was executed in 50% of all runs.
- The following branch to the end of the function (branch 1) was executed in 100% of all cases.
- The last branch to the end of the function (branch 2) was also executed in all runs.

# Concepts (1/2)

- The aim of **tuning** is the systematic performance increase, where the critical locations are identified and optimised.
- TIME determines the **runtime** of a process and thus the global effect of optimisations.
- The flat profile of GPROF shows how the runtime is divided between the individual functions.



## Concepts (2/2)

- The **structured** profile of GPROF shows how the runtime of a function is divided between sub-functions it has called.
- Using GCOV, it is possible to establish **how often** individual lines and branches were executed.
- The GPROF runtime profiling uses **sampling**, and the measured times are subject to statistical **inaccuracy**.
- In contrast, the number of calls and executions reported by GPROF and GCOV are **accurate**.